

▼ SOUMYA MUKHERJEE

CH24M571

M.Tech in INDUSTRIAL AI

Assignment - 2

pls find the notebook at : <https://github.com/aymuos/masters-practise-repo/tree/main/TERM4/Assignment2.ipynb>

```
1 %pip install numpy scikit-learn pandas tabulate matplotlib statsmodels
2

Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (2.0.2)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.12/dist-packages (1.6.1)
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packages (2.2.2)
Requirement already satisfied: tabulate in /usr/local/lib/python3.12/dist-packages (0.9.0)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (3.10.0)
Requirement already satisfied: statsmodels in /usr/local/lib/python3.12/dist-packages (0.14.5)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.16.0)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.5.0)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (3.1.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.3.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (4.22.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.3.1)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (25.0)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (3.2.2)
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.12/dist-packages (from statsmodels) (1.0.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->)
```

▼ Assignment 1

(Q1)

Approach used:

Alternating Location Allocation heuristics.

Breaks the problem into 2 alternating steps till the problem converges.

Phase - A : Allocation.

Warehouse location are fixed. We assign each dealership to the nearest warehouse based on Euclidian distance to minimise cost.

Assign dealership into warehouse j

if $d_{ij} < d_{ik}$ for all other warehouses

K.

Phase - B : Location Step.

Here, the assignment clusters are fixed. We want to find the optimal location of the warehouse within its assigned cluster using the centre of Gravity method which minimises the weighted sum of squared Euclidian distances

$$X_{\text{warehouse}} = \frac{\sum (x_i \cdot w_i)}{\sum w_i}$$

$$Y_{\text{warehouse}} = \frac{\sum (y_i \cdot w_i)}{\sum w_i}$$

The heuristics iterates b/won Phase A & Phase B.

→ we take a guess initial location

→ Allocate dealers to nearest guess

→ Relocate warehouse to COG of new cluster

→ Repeat until allocations no longer change (stabilized).

Thus, this method guarantees a local optima.

Approach used :

- ALA (Alternating Location-Allocation) Heuristic, which breaks the problem into two alternating steps until the solution stabilizes (converges).

Phase A: The Allocation Step

In this step, the locations of the warehouses are fixed. We assign each dealership to the nearest warehouse based on Euclidean distance to minimize cost.

Rule: Assign Dealership i to Warehouse j if $d_{ij} < d_{ik}$ for all other warehouses k . Theory: This creates a "Voronoi Partition" of the space, ensuring every customer is served by the closest facility.

Phase B: The Location Step (Center of Gravity)

In this step, the assignments (clusters) are fixed. We want to find the optimal location for the warehouse within its assigned cluster using the Center of Gravity (COG) method which minimizes the weighted sum of squared Euclidean distances. The coordinates (X, Y) for a warehouse are calculated as the weighted average of the coordinates of all dealerships assigned to it:

$$X_{warehouse} = \frac{\sum(x_i \cdot w_i)}{\sum w_i} \quad , \quad Y_{warehouse} = \frac{\sum(y_i \cdot w_i)}{\sum w_i}$$

Convergence of the algorithm :

The heuristic iterates between Phase A and Phase B.

- Guess initial locations.
- Allocate dealers to nearest guess.
- Re-locate the warehouse to the Center of Gravity of the new cluster.
- Repeat until the allocations no longer change (Stabilization).

This method guarantees a local optimum

```

1 import pandas as pd
2 import numpy as np
3 from scipy.spatial.distance import euclidean
4 from scipy.optimize import linprog

1 data = {
2     'Dealership': ['D1', 'D2', 'D3', 'D4', 'D5', 'D6', 'D7', 'D8'],
3     'x': [10, 20, 90, 85, 60, 65, 15, 20],
4     'y': [50, 45, 10, 15, 70, 75, 25, 30],
5     'Demand': [200, 150, 180, 170, 160, 140, 130, 120]
6 }
7 df = pd.DataFrame(data)

1 # 2. Function to Calculate Total Weighted Cost and Allocate
2
3 def calculate_total_cost(df, W):
4
5     """Calculates the total weighted cost and performs the allocation step."""
6
7
8     W1, W2 = W
9
10    # Calculate Euclidean distances to current warehouse locations
11    df['dist_W1'] = df.apply(lambda row: euclidean((row['x'], row['y']), W1), axis=1)
12    df['dist_W2'] = df.apply(lambda row: euclidean((row['x'], row['y']), W2), axis=1)
13
14    # Allocation: Assign to the closest warehouse
15    df['Assigned_W'] = np.where(df['dist_W1'] <= df['dist_W2'], 1, 2)
16    df['Min_Dist'] = np.minimum(df['dist_W1'], df['dist_W2'])
17    df['Weighted_Cost'] = df['Demand'] * df['Min_Dist']
18
19    total_cost = df['Weighted_Cost'].sum()
20    return total_cost, df
21
22
23 # 3. Center of Gravity (COG) Method
24 def center_of_gravity(cluster_df):
25
26     """Finds the Center of Gravity (COG) for a fixed cluster."""
27
28     if cluster_df.empty:
29         return np.array([0.0, 0.0])
30
31     total_demand = cluster_df['Demand'].sum()
32
33     x_cog = (cluster_df['x'] * cluster_df['Demand']).sum() / total_demand
34     y_cog = (cluster_df['y'] * cluster_df['Demand']).sum() / total_demand

```

```

35
36     return np.array([x_cog, y_cog])
37
38 #MAIN LOOP
39
40 # 4. Alternating Location-Allocation (ALA) Heuristic Main Loop
41 def ala_heuristic_cog(df, initial_W):
42     W = initial_W
43     prev_allocation = None
44
45
46     for iteration in range(1, 10):
47
48         # --- Allocation
49         cost, current_df = calculate_total_cost(df.copy(), W)
50         current_allocation = current_df['Assigned_W'].tolist()
51
52         # Check for convergence (Allocation stabilization)
53         if current_allocation == prev_allocation:
54             # Re-run the COG calculation one last time with the final allocation for precise coordinates
55             cluster1_df = current_df[current_df['Assigned_W'] == 1]
56             cluster2_df = current_df[current_df['Assigned_W'] == 2]
57             W1_new = center_of_gravity(cluster1_df)
58             W2_new = center_of_gravity(cluster2_df)
59             W = (W1_new, W2_new)
60             final_cost, final_df = calculate_total_cost(df.copy(), W)
61             return W, final_cost, final_df
62
63         prev_allocation = current_allocation
64
65         # --- Location Step (COG)
66
67         cluster1_df = current_df[current_df['Assigned_W'] == 1]
68         W1_new = center_of_gravity(cluster1_df)
69
70         cluster2_df = current_df[current_df['Assigned_W'] == 2]
71         W2_new = center_of_gravity(cluster2_df)
72
73         W = (W1_new, W2_new)
74
75     return W, cost, current_df

```

```

1
2
3 W_initial = (    np.array([16.75, 39.50]),
4                 np.array([75.31, 40.08]))
5
6 optimal_W, min_cost, final_df = ala_heuristic_cog(df.copy(), W_initial)
7
8
9 W1_final = optimal_W[0]
10 W2_final = optimal_W[1]
11 output_df = final_df[['Dealership', 'x', 'y', 'Demand', 'Assigned_W', 'Min_Dist', 'Weighted_Cost']].copy()
12 output_df.rename(columns={'Assigned_W': 'Warehouse', 'Min_Dist': 'Distance (km)', 'Weighted_Cost': 'Weighted Cost (units*km)'})
13
14 output_df['Warehouse'] = 'W' + output_df['Warehouse'].astype(str)
15 output_df['Distance (km)'] = output_df['Distance (km)'].round(2)
16 output_df['Weighted Cost (units*km)'] = output_df['Weighted Cost (units*km)'].round(2)
17
18 print(f"\nOptimal Warehouse 1 Location (W1): ({W1_final[0]:.2f}, {W1_final[1]:.2f}) km")
19
20 print(f"Optimal Warehouse 2 Location (W2): ({W2_final[0]:.2f}, {W2_final[1]:.2f}) km")
21
22 print(f"Minimum Total Transportation Cost: {min_cost:.2f} units * km")
23
24 print("-----")
25 print("\nFinal Dealership Assignments and Costs (Table View):")
26 print(output_df.to_markdown(index=False, floatfmt=".2f"))

```

Optimal Warehouse 1 Location (W1): (15.58, 39.33) km
Optimal Warehouse 2 Location (W2): (75.92, 40.08) km
Minimum Total Transportation Cost: 27646.92 units * km

Final Dealership Assignments and Costs (Table View):

Dealership	x	y	Demand	Warehouse	Distance (km)	Weighted Cost (units*km)
D1	10	50	200	W1	12.04	2407.92
D2	20	45	150	W1	7.18	1077.69
D3	90	10	180	W2	33.21	5977.47
D4	85	15	170	W2	26.67	4533.75
D5	60	70	160	W2	33.90	5423.35


D6	65	75	140	W2	36.59	5122.80
D7	15	25	130	W1	14.35	1864.88
D8	20	30	120	W1	10.33	1239.07

```

1 import matplotlib.pyplot as plt
2
3 def plot_warehouse_allocation(df, warehouses):
4
5     W1, W2 = warehouses
6
7     plt.figure(figsize=(12, 8))
8
9     colors = {1: 'blue', 2: 'green'}
10
11    # 1. Plot the Warehouses
12
13    # Warehouse 1
14    plt.scatter(W1[0], W1[1], c='blue', marker='*', s=400,
15                label='Optimal W1', edgecolors='black', zorder=10)
16    plt.text(W1[0], W1[1]+2, 'W1', fontweight='bold', ha='center', color='blue')
17
18    # Warehouse 2
19    plt.scatter(W2[0], W2[1], c='green', marker='*', s=400,
20                label='Optimal W2', edgecolors='black', zorder=10)
21    plt.text(W2[0], W2[1]+2, 'W2', fontweight='bold', ha='center', color='green')
22
23    # 2. Plot Dealerships and Connections
24    for index, row in df.iterrows():
25        w_id = int(row['Assigned_W'])
26        color = colors[w_id]
27        warehouse_coords = W1 if w_id == 1 else W2
28
29        # A. Draw Connection Line (Spider line)
30        plt.plot([row['x'], warehouse_coords[0]], [row['y'], warehouse_coords[1]],
31                  c=color, linestyle='--', alpha=0.3)
32
33        # B. Plot Dealership
34        plt.scatter(row['x'], row['y'], c=color, s=row['Demand'],
35                    alpha=0.6, edgecolors='black', label=f'Cluster {w_id}' if index in [0, 4] else "")
36
37        # C. Label Dealership
38        plt.text(row['x']+2, row['y'], row['Dealership'], fontsize=9)
39
40    # 3. Formatting the Graph
41    plt.title(f'Optimized Warehouse Locations & Allocations\n(Total Cost: {min_cost:.2f})', fontsize=14)
42    plt.xlabel('X Coordinate (Distance)')
43    plt.ylabel('Y Coordinate (Distance)')
44    plt.legend(loc='upper right')
45    plt.grid(True, linestyle=':', alpha=0.6)
46
47    plt.show()
48
49
50 plot_warehouse_allocation(final_df, optimal_W)

```

Optimized Warehouse Locations & Allocations (Total Cost: 27646.92)



Answer 1

Optimal Warehouse 1 Location (W1): (15.58, 39.33) km

Optimal Warehouse 2 Location (W2): (75.92, 40.08) km

Minimum Total Transportation Cost: 27646.92 units * km

Assignment 2

Question

An e-commerce company is planning to optimize its delivery fleet scheduling. The company has three distribution centers (DCs) and five delivery zones (DZs). Each DC has a fixed capacity in terms of the number of packages it can process daily, while each DZ has a fixed demand of packages to be delivered. The company wants to optimize the number of trucks to send from each DC to the delivery zones such that the total cost is minimized. The cost is determined by the distance between the DCs and DZs, and the goal is to ensure that each delivery zone's demand is met while not exceeding the capacity of the distribution centers

```

1 # given stuffs
2
3 DCs = ['DC1', 'DC2', 'DC3']
4 DZs = ['DZ1', 'DZ2', 'DZ3', 'DZ4', 'DZ5']
5 num_routes = len(DCs) * len(DZs)
6
7 # Cost Matrix is Flattened into a 1D cost vector
8 costs_matrix = np.array([
9     [4, 3, 6, 8, 9],
10    [5, 8, 3, 7, 6],
11    [9, 4, 7, 5, 2]
12 ])
13
14 c = costs_matrix.flatten()
15
16 # Capacities (b_ub): RHS of inequality constraints (DC Capacity <= 500, 600, 400)

```

```

17 b_ub = np.array([500, 600, 400])
18
19 # Demands (b_eq): RHS of equality constraints (DZ Demand == 200, 300, 350, 400, 250)
20 b_eq = np.array([200, 300, 350, 400, 250])
21

```

```

1 A_ub = np.zeros((len(DCs), num_routes))
2 for i, dc in enumerate(DCs):
3     # Sets 1s for the 5 variables corresponding to flow out of DC i
4     A_ub[i, i * len(DZs) : (i + 1) * len(DZs)] = 1
5
6 # A_eq: Demand Constraints (5 rows, 15 columns)
7 # Each row ensures: Sum(X_1j + X_2j + X_3j) == Demand_j
8
9
10 A_eq = np.zeros((len(DZs), num_routes))
11 for j, dz in enumerate(DZs):
12     # Sets 1s for the 3 variables corresponding to flow into DZ j (X_1j, X_2j, X_3j)
13     A_eq[j, j:len(DZs)] = 1
14
15 # Solve the Linear Program using linprog from scipy using highs method
16
17 result = linprog(c,
18                    A_ub=A_ub, b_ub=b_ub,
19                    A_eq=A_eq, b_eq=b_eq,
20                    method='highs')

```

```

1
2
3 if result.success:
4     optimal_flow = result.x.round(0)
5     optimal_cost = result.fun
6
7     # Reshape the 1D flow array back into a 3x5 matrix
8     flow_matrix = optimal_flow.reshape(len(DCs), len(DZs))
9     flow_df = pd.DataFrame(flow_matrix, index=DCs, columns=DZs)
10    flow_df['Total Dispatched'] = flow_df.sum(axis=1)
11
12    # Detailed assignment table
13    assignment_results = []
14    for i, dc in enumerate(DCs):
15        for j, dz in enumerate(DZs):
16            flow = flow_matrix[i, j]
17            if flow > 0:
18                unit_cost = costs_matrix[i, j]
19                assignment_results.append({
20                    'From_DC': dc,
21                    'To_DZ': dz,
22                    'Packages_Sent': int(flow),
23                    'Unit_Cost': unit_cost,
24                    'Total_Cost': flow * unit_cost
25                })
26    assignment_df = pd.DataFrame(assignment_results)
27
28    print(f"\nOptimal Solution Status: {result.message}")
29    print(f"Minimum Total Transportation Cost: {optimal_cost:.2f}")
30
31    print("\nOptimal Flow (Packages Sent from DC to DZ):")
32    print(flow_df.to_markdown(floatfmt=".0f"))
33
34    print("\nDetailed Optimal Assignment:")
35    print(assignment_df[['From_DC', 'To_DZ', 'Packages_Sent', 'Unit_Cost', 'Total_Cost']].to_markdown(index=
36
37 else:
38    print(f"Optimization failed. Status: {result.message}")

```

Optimal Solution Status: Optimization terminated successfully. (HiGHS Status 7: Optimal)
 Minimum Total Transportation Cost: 5,750.00

Optimal Flow (Packages Sent from DC to DZ):

	DZ1	DZ2	DZ3	DZ4	DZ5	Total Dispatched
DC1	200	300	0	0	0	500
DC2	0	0	350	250	0	600
DC3	0	0	0	150	250	400

Detailed Optimal Assignment:

From_DC	To_DZ	Packages_Sent	Unit_Cost	Total_Cost
DC1	DZ1	200	4	800.00
DC1	DZ2	300	3	900.00

DC2	DZ3		350		3		1,050.00	
DC2	DZ4		250		7		1,750.00	
DC3	DZ4		150		5		750.00	
DC3	DZ5		250		2		500.00	

Question 3

▼ QUestion

SportWear Inc. is a retail company specializing in athletic footwear. They want to forecast their monthly sales for the next year.

Historical data shows clear seasonal patterns with peaks during back-to-school season (August) and holiday season (December), as well as a general upward trend over the years. Given the following monthly sales data for the past 3 years (2021-2023) in thousands of units:

Soumya Mukherjee

(Q3)

(S1)

Components of the ETS (Error trend seasonality) model.

Trend: T_t is linearly upward.

Sales increases Y-o-Y.

Jan sales across years: $45 \rightarrow 50 \rightarrow 56 \dots$
 Dec sales : $80 \rightarrow 87 \rightarrow 94$ } in thousand units.

Seasonality (s_t): 12 months pattern is

Aug to dec is peak season

Feb is the lowest month.

Error: Random noise

In additive model forecast

$$\Rightarrow y_t = \text{Base level} + \text{Trend}_t + \text{Seasonality}_t + \text{Error}_t.$$

Yearly averages:

2021 Avg : 58.75

2022 Avg : 64.33

2023 Avg : 70.33

Avg growth = 5.8 units/year

$$\text{Avg growth} = T_t = 55.63$$

$$\text{Thus, trend equation} = T_t = 55.63 + 0.48t$$

where t = no of months

```

1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4

```

```

5
6 # Given DATA
7
8 data = {
9     'Month_Name': ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec'] * 3,
10    'Year': [2021]*12 + [2022]*12 + [2023]*12,
11    'Sales': [
12        45, 42, 48, 52, 55, 58, 62, 75, 65, 60, 63, 80, # 2021
13        50, 47, 53, 57, 60, 64, 68, 82, 70, 65, 69, 87, # 2022
14        56, 52, 58, 63, 66, 70, 74, 89, 76, 71, 75, 94 # 2023
15    ]
16 }
17
18 df = pd.DataFrame(data)
19
20
21 df['t'] = np.arange(1, len(df) + 1)
22
23
24 # DETERMINE TREND (using Linear Regression)
25
26 # Sales = Base + Slope * t
27 slope, intercept = np.polyfit(df['t'], df['Sales'], 1)
28
29 df['Trend_Component'] = intercept + slope * df['t']
30
31 print(f"Trend Equation: Sales = {intercept:.2f} + {slope:.2f} * t")
32
33
34 # 3. CALCULATE SEASONAL INDICES (Additive)
35
36 # De-trend the data (Actual - Trend)
37 df['Detrended'] = df['Sales'] - df['Trend_Component']
38
39 # Average the detrended values by month to get indices
40 seasonal_indices = df.groupby('Month_Name')['Detrended'].mean()
41
42 # Reorder indices to be Jan-Dec standard
43 month_order = ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec']
44 seasonal_indices = seasonal_indices.reindex(month_order)
45
46 print("\nSeasonal Indices:")
47 print(seasonal_indices.round(2))
48
49 # Map indices back to the main dataframe
50 df['Seasonal_Component'] = df['Month_Name'].map(seasonal_indices)
51
52
53 # 4. CALCULATE FITTED VALUES & MAPE
54
55 # Fitted Value = Trend + Seasonal
56 df['Fitted'] = df['Trend_Component'] + df['Seasonal_Component']
57
58 # Calculate Absolute Percentage Error for every row
59 df['APE'] = abs((df['Sales'] - df['Fitted']) / df['Sales']) * 100
60
61 # MAPE (Mean Absolute Percentage Error)
62 mape = df['APE'].mean()
63 print(f"\nModel MAPE: {mape:.2f}%")
64
65
66 # 5. FORECAST JAN - MAR 2024
67 future_months = pd.DataFrame({
68     'Month_Name': ['Jan', 'Feb', 'Mar'],
69     't': [37, 38, 39]
70 })
71
72 # Apply Forecast Formula
73 future_months['Trend_Forecast'] = intercept + slope * future_months['t']
74 future_months['Seasonal_Forecast'] = future_months['Month_Name'].map(seasonal_indices)
75 future_months['Final_Forecast'] = future_months['Trend_Forecast'] + future_months['Seasonal_Forecast']
76
77 print("\nForecast for Q1 2024:")
78 print(future_months[['Month_Name', 'Final_Forecast']].round(2))
79
80
81 # PLOTTING
82
83 plt.figure(figsize=(10, 6))
84 plt.plot(df['t'], df['Sales'], label='Actual Sales', marker='o')
85 plt.plot(df['t'], df['Fitted'], label='Model Fit', linestyle='--')
86 plt.plot(future_months['t'], future_months['Final_Forecast'], label='Forecast', color='red', marker='x')

```

```

87 plt.title('SportWear Inc: Sales vs ETS Model')
88 plt.xlabel('Month (t)')
89 plt.ylabel('Sales (000s)')
90 plt.legend()
91 plt.grid(True)
92 plt.show()

```

Trend Equation: Sales = 50.66 + 0.75 * t

Seasonal Indices:


Month_Name	Seasonal_Index
Jan	-10.03
Feb	-14.11
Mar	-8.86
Apr	-5.27
May	-3.02
Jun	-0.10
Jul	3.15
Aug	16.41
Sep	3.99
Oct	-1.75
Nov	1.17
Dec	18.42

Name: Detrended, dtype: float64

Model MAPE: 3.59%

Forecast for Q1 2024:

Month_Name	Final_Forecast
Jan	68.25
Feb	64.92
Mar	70.92



▼ Question 4

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 from statsmodels.tsa.stattools import acf, pacf
5 from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
6 from statsmodels.tsa.statespace.sarimax import SARIMAX
7 from sklearn.metrics import mean_squared_error
8
9 # given stuff
10 np.random.seed(42) # For reproducibility
11
12 # Parameters
13 n_days = 180
14 phi = 0.7      # AR parameter
15 theta = 0.3    # MA parameter
16 mu = 520       # Mean
17 sigma = 85     # Standard deviation
18

```


```

19 # Generate time index
20 t = np.arange(1, n_days + 1)
21
22 # Generate random error terms
23 e = sigma * np.random.randn(n_days)
24
25 # Add weekly seasonality (Sine wave with period 7)
26 seasonal_pattern = 50 * np.sin(2 * np.pi * t / 7)
27
28 # Generate ARIMA process (ARMA(1,1) + Seasonality)
29 y = np.zeros(n_days)
30 y[0] = mu + e[0]
31
32 for i in range(1, n_days):
33     # y[i] = mu + phi*(y[i-1] - mu) + e[i] + theta*e[i-1]
34     y[i] = mu + phi * (y[i-1] - mu) + e[i] + theta * e[i-1]
35
36 # Combine ARMA process with Seasonal pattern
37 daily_orders = y + seasonal_pattern
38
39 # Create a DataFrame
40 df = pd.DataFrame({'Day': t, 'Orders': daily_orders})
41 train = df.iloc[:170]
42 test = df.iloc[170:]
43
44 # --- 2. Visualization & Analysis (Task 1) ---
45
46 plt.figure(figsize=(12, 6))
47 plt.plot(df['Day'], df['Orders'], label='Daily Orders')
48 plt.title('FreshFood Daily Order Volumes (180 Days)')
49 plt.xlabel('Day')
50 plt.ylabel('Orders')
51 plt.legend()
52 plt.grid(True)
53 plt.show()
54
55
56
57 # Plot ACF and PACF to determine p, d, q
58 fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 8))
59 plot_acf(df['Orders'], lags=20, ax=ax1)
60 plot_pacf(df['Orders'], lags=20, ax=ax2)
61 plt.tight_layout()
62 plt.show()
63
64
65
66
67 # --- 3. Model Fitting & Forecasting (Task 2) ---
68
69 # Rationale for parameters:
70 # d=0: The data oscillates around a mean (520), so it is stationary.
71 # Seasonality: The problem states "Weekly seasonality" and lag 7 autocorrelation.
72 # Standard ARIMA(p,d,q) cannot handle seasonality well. We must use SARIMA.
73 # Based on the generation code (phi=0.7, theta=0.3), we expect AR=1, MA=1.
74 # Seasonal Order: We will add a seasonal component (1, 0, 1, 7) to capture the weekly pattern.
75
76 # Fit the model on the TRAINING set (First 170 days)
77 # Order = (p,d,q), Seasonal_Order = (P,D,Q,s)
78 model = SARIMAX(train['Orders'],
79                  order=(1, 0, 1),
80                  seasonal_order=(1, 0, 1, 7),
81                  trend='c') # 'c' adds a constant (intercept) close to the mean
82
83 model_fit = model.fit(disp=False)
84
85 print(model_fit.summary())
86
87 # Forecasts for days 171-180
88 forecast_result = model_fit.get_forecast(steps=10)
89 forecast_values = forecast_result.predicted_mean
90 conf_int = forecast_result.conf_int()
91
92 # --- 4. Evaluation
93
94 # Compare Forecasts with Actuals
95 comparison = pd.DataFrame({
96     'Day': test['Day'].values,
97     'Actual': test['Orders'].values,
98     'Forecast': forecast_values.values
99 })
100

```

```
101 print("\nForecast vs Actuals (Days 171-180):")
102 print(comparison)
103
104 # Calculate RMSE
105 rmse = np.sqrt(mean_squared_error(test['Orders'], forecast_values))
106 print(f"\nRoot Mean Square Error (RMSE): {rmse:.2f}")
107
108 # Plotting the Forecast
109 plt.figure(figsize=(12, 6))
110 plt.plot(train['Day'][-30:], train['Orders'][-30:], label='Training Data (Last 30 days)')
111 plt.plot(test['Day'], test['Orders'], label='Actual Values', color='green', marker='o')
112 plt.plot(test['Day'], forecast_values, label='Forecast', color='red', linestyle='--', marker='x')
113 plt.fill_between(test['Day'], conf_int.iloc[:, 0], conf_int.iloc[:, 1], color='pink', alpha=0.3, label='95%')
114 plt.title('FreshFood Delivery: Forecast vs Actual (Days 171-180)')
115 plt.xlabel('Day')
116 plt.ylabel('Order Volume')
117 plt.legend()
118 plt.grid(True)
119 plt.show()
120
121
```



Answers



▼ INterpretation

SARIMAX was used as the ARIMA Model due to the presence of seasonality. RMSE is 149.89, which is significantly higher than the standard deviation of the noise ($\sigma = 85$). The RMSE level ($\approx 85 - 90$) An RMSE of ~150 implies the model is missing a large, predictable pattern in the data.

```
1 # trying with ARIMAX but with sine wave as a feature (an exogenous variable)
2
3 from statsmodels.tsa.arima.model import ARIMA
4
5 # --- 1. Feature Engineering (Fourier Terms) ---
6
7 train['sin_wave'] = np.sin(2 * np.pi * train['Day'] / 7)
8 test['sin_wave'] = np.sin(2 * np.pi * test['Day'] / 7)
9
10 # --- 2. Fit ARIMAX Model ---
11 # standard ARIMA (p=1, d=0, q=1) but add the 'exog' (exogenous) variable.
12
13 model = ARIMA(endog=train['Orders'],
14                 exog=train[['sin_wave']],
15                 order=(1, 0, 1),
16                 trend='c') # Constant intercept
17
18 model_fit = model.fit()
19
```



```
warn('Non-stationary starting seasonal autoregressive'
      SARIMAX Results
```

```
=====
=====
Dep. Variable: Orders No. Observations: 170
Model: SARIMAX(1, 0, 1)x(1, 0, 1, 7) Log Likelihood -991.513
Date: Sun, 07 Dec 2025 AIC 1995.026
Time: 22:13:32 BIC 2013.841
Sample: 0 HQIC 2002.661
- 170
Covariance Type: opg
=====
=====
```

	coef	std err	z	P> z	[0.025	0.975]
intercept	155.5249	149.417	1.041	0.298	-137.327	448.376
ar.L1	0.5659	0.095	5.970	0.000	0.380	0.752
ma.L1	0.4020	0.102	3.957	0.000	0.203	0.601
ar.S.L7	0.2953	0.624	0.474	0.636	-0.927	1.518
ma.S.L7	-0.1546	0.635	-0.243	0.808	-1.400	1.090
sigma2	6768.5162	767.794	8.816	0.000	5263.668	8273.364

```
=====
=====
```


Ljung-Box (L1) (Q):	0.00	Jarque-Bera (JB):	0.16
Prob(Q):	0.98	Prob(JB):	0.93
Heteroskedasticity (H):	1.15	Skew:	-0.02
Prob(H) (two-sided):	0.59	Kurtosis:	2.86

```
=====
=====
```

...

8 179 636.911558 518.356935
 9 180 792.743256 515.120262

Root Mean Square Error (RMSE): 149.89



SARIMAX Results

```

=====
=====
Dep. Variable:          Orders  No. Observations:      170
Model:      ARIMA(1, 0, 1)  Log Likelihood: -985.245
Date:      Sun, 07 Dec 2025  AIC:            1980.490
Time:      23:36:11    BIC:            1996.169
Sample:     0   HQIC:           1986.852
           - 170
Covariance Type: opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
const	502.5826	22.039	22.804	0.000	459.387	545.778
sin_wave	56.7231	13.529	4.193	0.000	30.208	83.239
ar.L1	0.6473	0.083	7.768	0.000	0.484	0.811
ma.L1	0.2741	0.110	2.495	0.013	0.059	0.489
sigma2	6299.9005	730.955	8.619	0.000	4867.254	7732.547

```

=====
Ljung-Box (L1) (Q):      0.01  Jarque-Bera (JB):       0.47
Prob(Q):                0.94  Prob(JB):                  0.79
Heteroskedasticity (H):  1.12  Skew:                     0.06
Prob(H) (two-sided):    0.66  Kurtosis:                 2.78
```

=====

=====

Warnings:

...

8 179 636.911558 479.654598
9 180 792.743256 448.371144

New Root Mean Square Error (RMSE): 150.24

